

Name: Chitradevi Maruthavanan

Week -6 Video Activity

Activity 11.1 –Storage and Indexes

1. If you are in psql, run \d books to see that there is an index on books.id

Answer:

Yes, I can see that there is an index on books.id

```
su22adb20=> \d books
              Table "su22adb20.books"
  Column      | Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
 id           | integer   |           | not null |
 title        | text      |           |          |
 pagecount    | integer   |           |          |
 genre        | text      |           |          |
 authorid     | integer   |           |          |
 pubid        | integer   |           |          |
Indexes:
    "books_pkey" PRIMARY KEY, btree (id)
    "books_title_idx" btree (title)
```

2. Create an index on books.pagecount

Answer:

```
CREATE INDEX ON books (pagecount);
```

3. Write a query that take advantage of that index.

Answer:

```
SELECT * FROM books WHERE pagecount =100;
```

Activity 11.2 – Index Matching

For the below – indicate if the index matches the predicate or not

Assume index on books(id)

1. id = 5132 **yes**
2. title = 'It' AND id = 5132 **yes**
3. title = 'It' OR id = 5132 **NO**

Assume index on books(genre,pagecount)

4. genre = 'Horror' AND pagecount > 2000 **Yes**
5. genre = 'Horror' **yes**
6. pagecount > 2000 **No**

Activity 12.1 –Index Structure

1. Explain SELECT * FROM agent WHERE id = 5;

Answer:

Yes, PostgreSQL uses an index

```
su22adb20=> Explain SELECT * FROM agent WHERE agent_id = 5;
               QUERY PLAN
-----
Index Scan using agent_pkey on agent  (cost=0.28..8.29 rows=1 width=54)
   Index Cond: (agent_id = 5)
(2 rows)
```

2. Explain SELECT * FROM agent WHERE id >5;

Answer:

No, PostgreSQL does not use an index

```
su22adb20=> Explain SELECT * FROM agent WHERE agent_id > 5;
               QUERY PLAN
-----
Seq Scan on agent  (cost=0.00..16.27 rows=658 width=54)
   Filter: (agent_id > 5)
(2 rows)
```

Activity 12.2 –Clustered and unclustered Indexes

1. Run: `SELECT * FROM books;`

Answer:

id	title	pagecount	genre	authorid	pubid
1	It	1138	Horror	10	100
2	Hamlet	500	Tragedy	13	103
3	I Know Why the Caged Bird Sings	304	Autobiographical	14	102
4	A Suitable Boy	1349	Drama/Romance	15	103
5	The Joy Luck Club	288	Drama	16	104
6	Like Water for Chocolate	256	Romance/Tragedy	17	105
7	Tita's Diary	294	Romance/Diary	17	
8	From Heaven Lake	464	Travel	15	102
9	Kite Runner	371	Historical/Drama	18	106
10	The Vanishing Half	352	Historical/Drama	19	106
11	September Love	224	Romance	20	107
12	The Nickel Boys	224	Historical	21	108
13	The Alchemist	163	Fantasy/Adventure	22	103
14	Love and Misadventure	176	Romance	20	107
15	The Authenticity Project	384	Romance	23	102
16	If I never met you	543	Romance	12	132

(16 rows)

Note:

Before cluster the relation is not sorted on any attribute (the relation is almost sorted on id)

2. **CREATE** an index on books on pagecount using the command

`CREATE INDEX ON books(pagecount);`

Answer:

```
su22adb20=> CREATE INDEX ON books(pagecount);  
CREATE INDEX
```

3. If you are using `psql`, you can use `\d books` to see the name of the new index

```
su22adb20=> \d books  
          Table "su22adb20.books"  
  Column      | Type      | Collation | Nullable | Default  
-----+-----+-----+-----+-----  
 id            | integer   |           | not null |  
 title         | text      |           |          |  
 pagecount     | integer   |           |          |  
 genre         | text      |           |          |  
 authorid      | integer   |           |          |  
 pubid         | integer   |           |          |  
Indexes:  
 "books_pkey" PRIMARY KEY, btree (id)  
 "books_pagecount_idx" btree (pagecount)  
 "books_pagecount_idx1" btree (pagecount)  
 "books_title_idx" btree (title)
```

4. Cluster books using the command:

```
CLUSTER books USING books_pagecount_index;
```

```
su22adb20=> CLUSTER books USING books_pagecount_idx;  
CLUSTER
```

5. Run: SELECT * FROM books; again

id	title	pagecount	genre	authorid	pubid
13	The Alchemist	163	Fantasy/Adventure	22	103
14	Love and Misadventure	176	Romance	20	107
12	The Nickel Boys	224	Historical	21	108
11	September Love	224	Romance	20	107
6	Like Water for Chocolate	256	Romance/Tragedy	17	105
5	The Joy Luck Club	288	Drama	16	104
7	Tita's Diary	294	Romance/Diary	17	
3	I Know Why the Caged Bird Sings	304	Autobiographical	14	102
10	The Vanishing Half	352	Historical/Drama	19	106
9	Kite Runner	371	Historical/Drama	18	106
15	The Authenticity Project	384	Romance	23	102
8	From Heaven Lake	464	Travel	15	102
2	Hamlet	500	Tragedy	13	103
16	If I never met you	543	Romance	12	132
1	It	1138	Horror	10	100
4	A Suitable Boy	1349	Drama/Romance	15	103
(16 rows)					

Note:

After cluster the relation is sorted on pagecount.

I cannot count on postgres giving me a sorted result like this. In the instance, postgres will return the tuples in the order they are stored; however, the database can return tuples in any order it wants to.